

The background of the slide is decorated with several realistic water droplets and bubbles of varying sizes. Some are perfectly spherical, while others are more elongated or flattened, suggesting they are on a surface. They are rendered with soft shadows and highlights, giving them a three-dimensional appearance. The droplets are scattered across the slide, with a cluster of three small ones near the top center, a large one on the left side, a very large one on the right side, and a small one at the bottom right.

Marketing Campaign Prediction

Reddy Charan Pulivendula


Pandya Harsh Maheshkumar

Sheta Rushank Ghanshyam



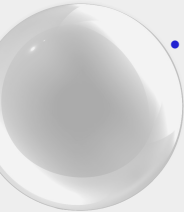
Abstract



- On average a company spends 4% to 25% of their net revenue into Marketing campaigns and Advertisements
 - While some companies like Salesforce and Tableau spent around 50% of their revenue in marketing campaigns
 - Targeted marketing is used to get good returns from this investment which aims at identifying groups of customers which are highly likely to become future customers
 - Customer segmentation allows us to group customers based on demographics, geography, interests, user activity, etc.
- 



Problem Statement

- Marketing campaign analysis allows companies to find customers which are highly likely to become future customer
 - Using historical data from previous marketing campaigns to predict the response of a new customer applying predictive modeling
 - Use customer demographics like age, income, marital status, etc. to classify if a new customer will respond positively or negatively to the campaigns
- 



Related Work

Performed steps like:

- EDA on Data Set
 - Correlation Analysis
 - Handling Missing Data
 - Feature Engineering
 - Binary Classification
- 


Data

	target	month	duration	age	gender	job	maritalStatus	education	creditFailure
0	0	may	166	30	female	worker	married	highSchool	no
1	0	oct	183	42	female	manager	married	uniGraduated	no
2	0	jun	227	26	female	services	single	highSchool	no
3	0	jun	31	34	male	unemployed	divorced	uniGraduated	yes
4	0	may	1231	48	male	worker	married	secondarySchool	no


accountBalance	house	credit	contactType	numberOfContacts	daySinceLastCampaign	numberOfContactsLastCampaign	lastCampaignResult
-202	no	no	unknown	2	NaN	0	unknown
2463	no	no	cellPhone	2	NaN	0	unknown
2158	yes	yes	landline	1	NaN	0	unknown
75	yes	no	unknown	3	NaN	0	unknown
559	yes	no	unknown	2	NaN	0	unknown

Univariate Analysis

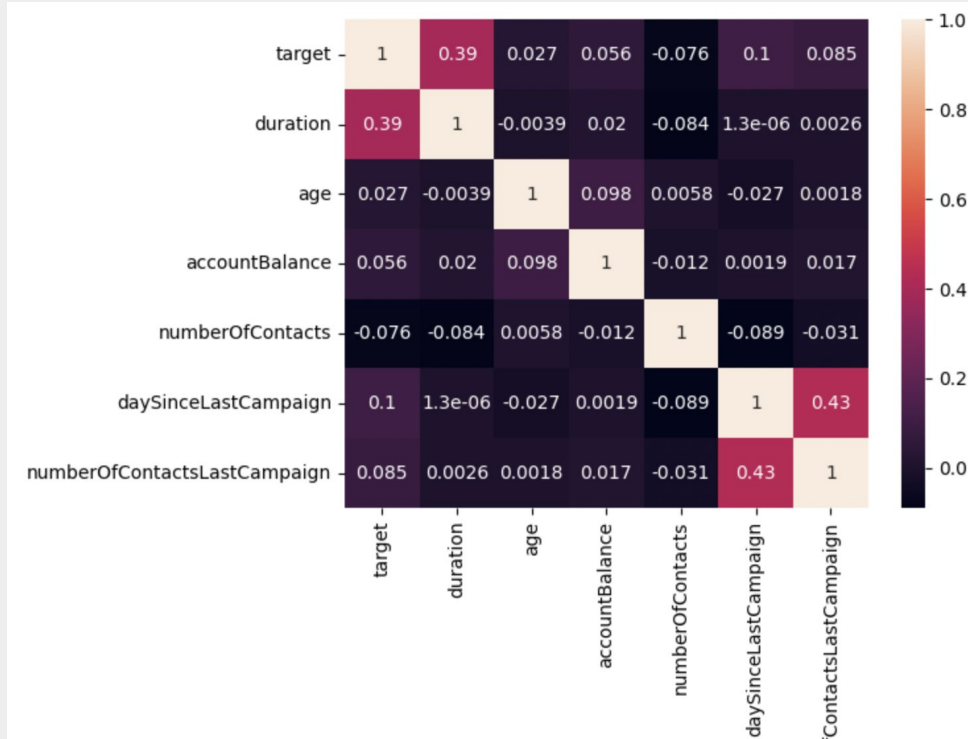
Campaign success ratio ratio during months of the year



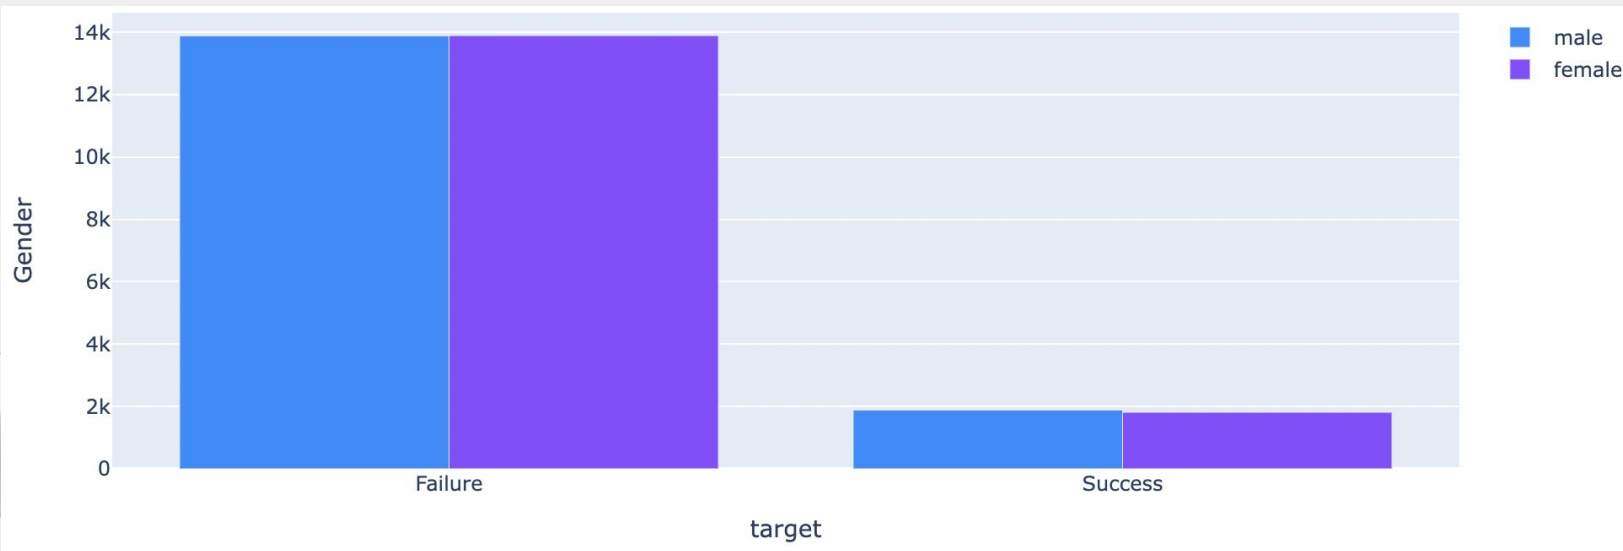
Univariate Analysis




Correlation Analysis



Bivariate Analysis



Bivariate Analysis



Handling null values

```
[14]: data.isnull().sum()
```

```
[14]: id                0
      target            0
      day              0
      month            0
      duration          0
      contactId        0
      age              0
      gender           0
      job              0
      maritalStatus    0
      education        0
      creditFailure    0
      accountBalance   0
      house            0
      credit           0
      contactType      0
      numberOfContacts 0
      daySinceLastCampaign 25742
      numberOfContactsLastCampaign 0
      lastCampaignResult 0
      dtype: int64
```

```
[15]: data['daySinceLastCampaign'].fillna(-1, inplace=True)
```

Replacing null values with -1, which implies that the customer did not respond positively in previous campaign

Data pre-processing

```
[ ]: # Encoding categorical attributes  
df = pd.get_dummies(final_data, columns = categorical_cols[1:])
```

```
[75]: categorical_cols[1:]
```

```
[75]: ['month',  
      'gender',  
      'job',  
      'maritalStatus',  
      'education',  
      'creditFailure',  
      'house',  
      'credit',  
      'contactType',  
      'lastCampaignResult']
```


Logistic Regression

```
[99]: from sklearn.linear_model import LogisticRegression
      clf = LogisticRegression(random_state=0).fit(X_train, y_train)
      from sklearn import metrics
      pred = clf.predict(X_test)
      print()
      print("Test Accuracy:", metrics.accuracy_score(y_test, pred))
      print()
      y_train_pred = clf.predict(X_train)
      train_acc = metrics.accuracy_score(y_train, y_train_pred)
      print("Train Accuracy:", train_acc)

      plot_confusion_matrix(clf, X_test, y_test)
```

Test Accuracy: 0.9037484116899619

Train Accuracy: 0.9005264113269196



Decision Tree

```
[102]: from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import plot_confusion_matrix

clf = DecisionTreeClassifier(random_state=0)
clf.fit(X_train,y_train)
predictions = clf.predict(X_test)

from sklearn import metrics
print()


# using metrics module for accuracy calculation
print("Test Accuracy: ", metrics.accuracy_score(y_test, predictions))

y_train_pred = clf.predict(X_train)
train_acc = metrics.accuracy_score(y_train, y_train_pred)
print("Train Accuracy:", train_acc)

plot_confusion_matrix(clf, X_test, y_test)
```

Test Accuracy: 0.8698644642100805

Train Accuracy: 1.0



Random Forest

```
[100]: from sklearn.ensemble import RandomForestClassifier
      clf = RandomForestClassifier(n_estimators = 200, bootstrap= True)

      # Training the model on the training dataset
      # fit function is used to train the model using the training sets as parameters
      clf.fit(X_train, y_train)

      # performing predictions on the test dataset
      y_pred = clf.predict(X_test)

      # metrics are used to find accuracy or error
      from sklearn import metrics
      print()


      # using metrics module for accuracy calculation
      print("Accuracy with 200 estimators : ", metrics.accuracy_score(y_test, y_pred))

      y_train_pred = clf.predict(X_train)
      train_acc = metrics.accuracy_score(y_train, y_train_pred)
      print("Train Accuracy:", train_acc)

      plot_confusion_matrix(clf, X_test, y_test)
```

Accuracy with 200 estimators : 0.9073485811096993

Train Accuracy: 1.0



Boosting

```
[101]: from sklearn.ensemble import AdaBoostClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import confusion_matrix

classifier = AdaBoostClassifier(
    DecisionTreeClassifier(max_depth = 1),
    n_estimators = 200
)
classifier.fit(X_train, y_train)
predictions = classifier.predict(X_test)


from sklearn import metrics
print()

# using metrics module for accuracy calculation
print("Accuracy: ", metrics.accuracy_score(y_test, predictions))
print()

y_train_pred = classifier.predict(X_train)
train_acc = metrics.accuracy_score(y_train, y_train_pred)
print("Train Accuracy:", train_acc)
plot_confusion_matrix(clf, X_test, y_test)
```

Accuracy: 0.904595510376959

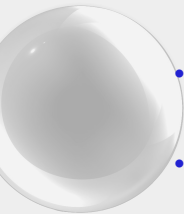
Train Accuracy: 0.9019332002178254





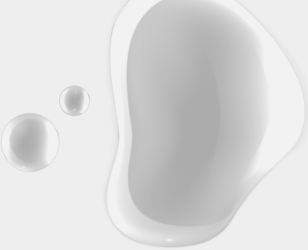
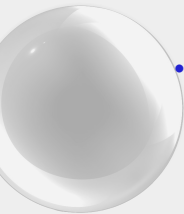
Results and Challenges



- Random Forest classifier results in best test accuracy
 - Models are highly training data specific
 - Data authenticity and Quality issues
 - Many irrelevant features
 - Limited availability of open source data
 - Class Imbalance
- 



Future Scope

- Apply methods to reduce impact of class imbalance on the classifier
 - Experiment with different combination of features
 - Optimize current models
 - Utilize more classification metrics to assess model's performance
- 
- 



References

- <https://www.kaggle.com/code/khanimar/bi-marketing-campaign-eda-analysis-prediction/data>
 - <https://towardsdatascience.com/how-to-predict-the-success-of-your-marketing-campaign-579fbb153a97>
 - <https://www.kaggle.com/code/seananguyen/marketing-campaign-analysis-python#3.-Data-Visualizations>
 - <https://waypointmc.com/blog/analyzing-marketing-results#:~:text=What%20is%20Marketing%20Analysis%3F,improve%20future%20conversions%20or%20sales.>
- 